

Implementation of Dec-POMDP Environment in Spann 2006

Baiting Zhu

Table of Contents

1. DEC-POMDP vs. POMDP
2. Assumptions of the experiment environment
3. Transition Function
4. Reward Function
5. Observation Function
6. Further Improvements

DEC-POMDP vs. POMDP

What makes a POMDP “decentralized”?

- More than 1 agent

agentList = ["agentOne", "agentTwo", ...]

- Agents taking actions simultaneously and are independent of each other

allAgentsActionTable = {"agentOne": (1, 0), "agentTwo": (0, 0), ...}


- Agents have (limited) observations

allAgentsObsTable = {"agentOne": "left", "agentTwo": "heaven-right", ...}

- Agents may communicate (share observations) at a cost
- Each agent only considers its local state + shared information about the environment

NOT local states of other agents

Spaan's Environment



Two agents start from the star and move independently to the left, right, up, down, or stay. Any movement has -0.1 cost.

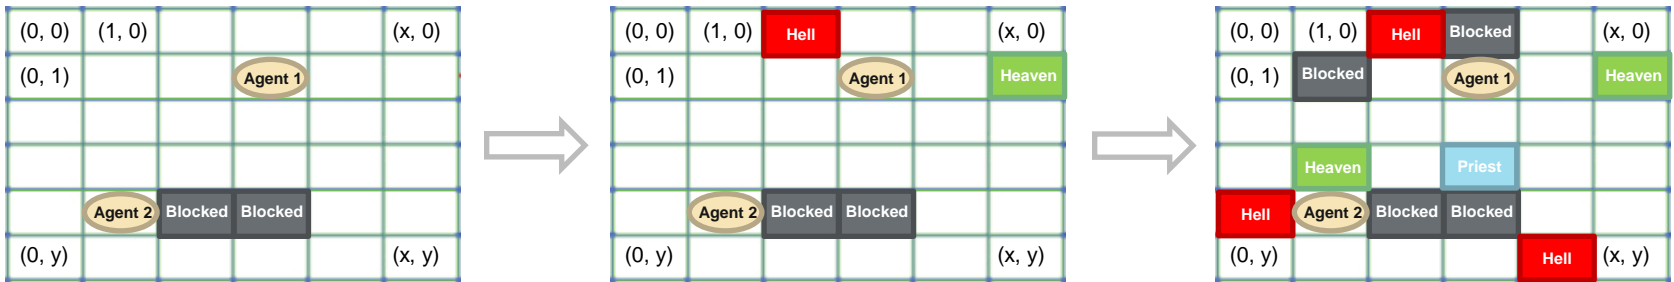
One of the bottom-left or bottom-right grid is Heaven, and the other one is Hell. This is unknown to the agents at the start.

Agents know which one is the Heaven when visiting the Priest. Doing so results in a -2 cost.

Once agents meet up in a grid that is Heaven or Hell, the game ends. With a $+10$ reward or -10 punishment.


Our Environment

- Any number of **n agents** may take actions simultaneously in a **2D rectangular** environment of size $x \times y$
- Any grid within the environment may be “**Blocked**” (unable to enter nor start from)
- Any number of “**Heaven**” and “**Hell**” may exist in unblocked grids, but they cannot be in the same grid. Each “Heaven” / “Hell” may also have unique **reward** / **punishment** values
- Any number of “**Priest**” may exist in unblocked grids that are not Heaven nor Hell. Each “Priest” may have unique cost



Game Movements

- Each agent must be in an unblocked grid in the map at any time stamp t
- Each agent may take actions $(1, 0)$ $(0, 1)$ $(-1, 0)$ $(0, -1)$ $(0, 0)$ at each time stamp t , correspondingly representing **Move Right, Move Down, Move Left, Move Up, and Stay**
- When an action is leading the agent **out of boundary** or **into a Blocked grid**, agent's action result in no movement



Block & Boundary are not reachable:

Agent 1, move left: location unchanged, due to block
Agent 1, move up: location unchanged, due to boundary

Heaven / Hell / Priest can be reached:

Agent 2, move up: location changes
Agent 2, move up: location changes
Agent 2, move up: location changes
Agent 2, move up: location unchanged

Game Reward and Punishment


$allAgentsNormalCost = \{ "agentOne": -0.1, "agentTwo": -0.2, ... \}$

- No matter the action an agent chooses to do (including stay), and no matter the consequence of the action, agent pays a **normal cost**. This cost may be different for each agent
- When all agents meet in **the same Heaven or Hell**, game ends. All agents take the award / punishment
- When game ends, further actions result in no cost nor movement
- Being in different Heaven or Hell doesn't lead to an end of the game
- *Other than normal cost and ending award / punishment, each agent may experience different reward /

punishment along the way

Agent 1 move right; Agent 2 moves up
-> Happy Ending

Agent 1 move down; Agent move left
-> Bad Ending




Game cannot end in one step, because agent cannot meet up in the SAME Heaven or Hell



Game Observations and Priest

- If an agent is not in a Priest grid, an agent observes whether there is a wall on its left and/or right. Note that this only relates to the horizontal (x) direction of the map, but not agent's orientation. May observe **“left”**, **“right”**, **“none”**, or **“both”**
- If an agent is in a Priest grid, it takes a cost and knows the location of each Heaven. Different priest may have different costs




Agent 1: observes “both”
Agent 2: observes “left”
Agent 3: observes “right”
Agent 4: observes “neither”

← $-x$ wall-obs direction $+x$ →

Agent 5: observes location of each Heaven

More Rules about Priest Cost

- If agent **stays in the Priest grid**, or its actions has **no resulting change in location** (into boundary or block), it **doesn't** take the Priest cost again
- Getting out from the Priest grid then **re-enters will result** in the same Priest cost
- Otherwise, the Priest grid should be all the same as a normal grid



Agent 1, move down: Priest cost paid (visiting)

Agent 2, move down: NO cost (blocked \equiv staying)

Agent 2, stays: NO cost (staying)

Agent 2, move left then right: Priest cost paid (re-entering \equiv visiting)

Transition Function (All Agents)

- $p(s'|s, a)$ where s', s, a are vectors representing the resulting states, original states, and actions of all agents

```
class TransitionFunction(object):

    def __init__(self, checkIfAllInHeaven, checkIfAllInHell, findSPRime):
        self.checkIfAllInHeaven = checkIfAllInHeaven
        self.checkIfAllInHell = checkIfAllInHell
        self.findSPRime = findSPRime

    def __call__(self, allAgentsPositionTable, allAgentsActionTable, allAgentsSPRimeTable):
        allAgentsInHeaven = self.checkIfAllInHeaven(allAgentsPositionTable, self.findSPRime.heavenDict)
        allAgentsInHell = self.checkIfAllInHell(allAgentsPositionTable, self.findSPRime.hellDict)
        if allAgentsInHeaven or allAgentsInHell:
            allSPRimeExpected = allAgentsPositionTable
        else:
            allSPRimeExpected = {agent: self.findSPRime(allAgentsPositionTable[agent], allAgentsActionTable[agent]) for agent in allAgentsActionTable.keys()}

        return int(allSPRimeExpected == allAgentsSPRimeTable)
```

Transition Function (One Agent)

- $p(s'|s, a)$ where s', s, a are vectors representing the resulting states, original states, and actions of all agents

```
def checkIfAllInHeaven(allAgentsPositionTable, heavenDict):
    if all(pos in heavenDict.keys() for pos in allAgentsPositionTable.values())\
        and len(set(allAgentsPositionTable.values())) == 1:
        return True
    else:
        return False

def checkIfAllInHell(allAgentsPositionTable, hellDict):
    if all(pos in hellDict.keys() for pos in allAgentsPositionTable.values())\
        and len(set(allAgentsPositionTable.values())) == 1:
        return True
    else:
        return False
```

```
class FindSPrime(object):
    def __init__(self, minX, minY, maxX, maxY, blockList, hellDict, heavenDict):
        self.minX = minX
        self.minY = minY
        self.maxX = maxX
        self.maxY = maxY
        self.blockList = blockList
        self.hellDict = hellDict
        self.heavenDict = heavenDict

    def __call__(self, s, action):
        x, y = s
        dx, dy = action
        sPrimeConsideringBoundary = (max(self.minX, min(x+dx, self.maxX)), max(self.minY, min(y+dy, self.maxY)))
        sPrime = s if sPrimeConsideringBoundary in self.blockList else sPrimeConsideringBoundary
        return sPrime
```

heavenDict = {(1, 2): 10, (2, 3): 9}

hellDict = {(1, 1): -5, (4, 3): -9.6}

hellDict = [(1, 0), (4, 6), (7, 2)]

allAgents[...].Table = {"agentOne": ..., "agentTwo": ...}

Reward Function (All Agents)

- $R(s, a, s')$ where s, a, s' are vectors representing the, original states, actions, and resulting states of all agents

```
class RewardFunction(object):  
    def __init__(self, allAgentsNormalCost, priestDict, hellDict, heavenDict, allAgentsMiddleStageDict):  
        self.allAgentsNormalCost = allAgentsNormalCost  
        self.allAgentsMiddleStageDict = allAgentsMiddleStageDict  
        self.priestDict = priestDict  
        self.hellDict = hellDict  
        self.heavenDict = heavenDict  
  
    def __call__(self, allAgentsPositionTable, allAgentsActionTable, allAgentsSPrimeTable):  
        allAgentsInHeaven = tt.checkIfAllInHeaven(allAgentsPositionTable, self.heavenDict)  
        allAgentsInHell = tt.checkIfAllInHell(allAgentsPositionTable, self.hellDict)  
        if not allAgentsInHeaven and not allAgentsInHell:  
            allAgentsInHeaven_after = tt.checkIfAllInHeaven(allAgentsSPrimeTable, self.heavenDict)  
            allAgentsInHell_after = tt.checkIfAllInHell(allAgentsSPrimeTable, self.hellDict)  
            findAgentReward = FindAgentReward(self.priestDict, self.hellDict, self.heavenDict, allAgentsInHeaven_after, allAgentsInHell_after)  
            rewardDict = {agent: findAgentReward(allAgentsPositionTable[agent], allAgentsSPrimeTable[agent], self.allAgentsNormalCost[agent], self.allAgentsMiddleStageDict[agent])\n                           for agent in allAgentsPositionTable.keys()}  
        else:  
            rewardDict = {agent: 0 for agent in allAgentsActionTable.keys()}  
        return rewardDict
```

Total reward = sum of reward of all agents \Longrightarrow e.g.: `sum(list(rewardDict.values()))`

Reward Function (One Agent)

- $R(s, a, s')$ where s, a, s' are vectors representing the, original states, actions, and resulting states of all agents

```
import DecPOMDP_transitionTable as tt

class FindAgentReward(object):

    def __init__(self, priestDict, hellDict, heavenDict, allAgentsInHeaven_after, allAgentsInHell_after):
        self.priestDict = priestDict
        self.hellDict = hellDict
        self.heavenDict = heavenDict
        self.allAgentsInHeaven_after = allAgentsInHeaven_after
        self.allAgentsInHell_after = allAgentsInHell_after

    def __call__(self, s, sPrime, normalCost, middleStageDict):
        if self.allAgentsInHeaven_after:
            return normalCost + self.heavenDict[sPrime]
        elif self.allAgentsInHell_after:
            return normalCost + self.hellDict[sPrime]
        else:
            if s == sPrime:
                return normalCost
            elif sPrime in self.priestDict.keys():
                return normalCost + self.priestDict[sPrime]
            else:
                return normalCost + middleStageDict[sPrime]
```

Observation Function (All Agents)

- $p(o|s, a)$ where o, s, a are vectors representing the, observations original states, and actions of all agents
- For this project, ignore the communication observations (future improvement)

```
class ObservationFunction(object):  
  
    def __init__(self, findOneAgentObsTable):  
        self.findOneAgentObsTable = findOneAgentObsTable  
  
    def __call__(self, allAgentsSPrimeTable, allAgentsActionTable, allAgentsObsTable):  
        allAgentsObsExpected = {agent: self.findOneAgentObsTable(allAgentsSPrimeTable, agent) for agent in allAgentsActionTable.keys()}  
        return int(allAgentsObsTable == allAgentsObsExpected)
```

Observation Function (One Agent)

- $p(o|s, a)$ where o, s, a are vectors representing the, observations original states, and actions of all agents
- For this project, ignore the communication observations (future improvement)

```
class FindOneAgentObsTable(object):  
  
    def __init__(self, minX, minY, maxX, maxY, blockList, priestDict, hellDict, heavenDict):  
        self.minX = minX  
        self.minY = minY  
        self.maxX = maxX  
        self.maxY = maxY  
        self.blockList = blockList  
        self.priestDict = priestDict  
        self.hellDict = hellDict  
        self.heavenDict = heavenDict  
  
    def __call__(self, allAgentsSPrimeTable, agent):  
        s = allAgentsSPrimeTable[agent]  
        if s in self.priestDict.keys():  
            if list(self.heavenDict.keys())[0] == (0, 3):  
                return 'heaven-Left'  
            else:  
                return 'heaven-right'  
        else:  
            x, y = s  
            leftIsWall = (x == self.minX or (x-1, y) in self.blockList)  
            rightIsWall = (x == self.maxX or (x+1, y) in self.blockList)  
            if leftIsWall and rightIsWall:  
                return 'both'  
            elif leftIsWall:  
                return 'Left'  
            elif rightIsWall:  
                return 'right'  
            else:  
                return 'neither'
```

Future Improvements

- Implement the communication component
- Add choices of observation function. Currently, to line up with set up in *Decentralized Planning under Uncertainty for Teams of Communicating Agents (Spaan et al.)*, the environment uses “heaven-left” and “heaven-right” to represent the location of Heaven and Hell. In scenarios of multiple Heavens and Hells, the functions need to be adjusted to show the correct outputs.
- Consider scenarios when users not giving the correct input. Throw Warnings or Errors



Thank You
